

EBU7240

Computer Vision

- Restoration: spatial filtering -

Semester 1, 2021

Changjae Oh



water



coffee



air



snapchat



aquarium



oil



instagram



icon



cartridge



sand



Noodlemans CubeCart Plu...
noodlemans.co.uk · In stock



CleverSpa Filters | J D ...
jdwilliams.co.uk · In stock



File:Filter.svg - Wikimedia C...
commons.wikimedia.org



Lay-Z-Spa Hot Tub Filter Ca...
amazon.co.uk



Hydraulic Filter Element G01369...
uk.rs-online.com · In stock



ÖVERST 3-piece metal filter ...
ikea.com · In stock



Hydraulic Filter, steering system MA...
autodoc.co.uk · In stock



Air Filter Induction Kit Sports ...
amazon.co.uk



Gal.) & Larger Wet Dry Vacu...
homedepot.ca · In stock



CleverSpa® Water Filter Pack
therange.co.uk · In stock



90mm High Flow Pleated C...
ramair-filters.co.uk · Out of st...



AKORD Microphone Sw...
amazon.co.uk



Best Home Water Filters | W...
popularmechanics.com



Filter (band) - Wikipedia
en.wikipedia.org

Content

- Neighborhood in an image ✓
- Convolution: review
 - From 'Signal Processing' Lecture ✓
- Spatial Filtering
 - Low-pass (or high-pass) filter ✓
 - Gaussian (or Laplacian) filter
 - Mean, median, or mode filter
 - Advanced filters: Bilateral filter, Non-local means filter

Neighborhood in an image

- Spatial filtering is performed for each pixel (i, j) using neighborhoods of (i, j) .
 - $I(i, j)$: input image $O(i, j)$: output image
 - $w(s, t)$: filtering kernel (a.k.a. mask)

$$O(i, j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t)I(i + s, j + t)$$

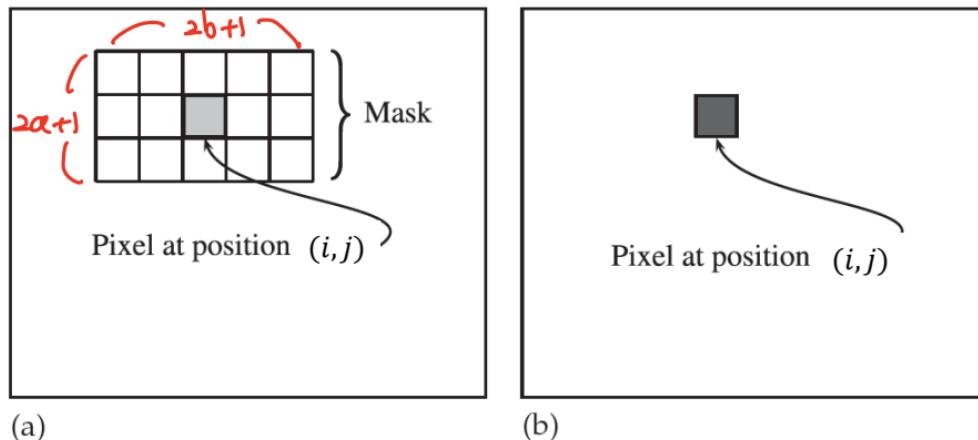


- The filtering operation is defined according to what kind of $w(s, t)$ is used.
- The filter serves as an essential building block for many applications.
 - Blurring, sharpening, image restoration, and so on

Spatial Filtering: Filter Kernel (Mask)

- Filter kernel should be defined accordingly, depending on applications.

$$O(i, j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t)I(i + s, j + t)$$

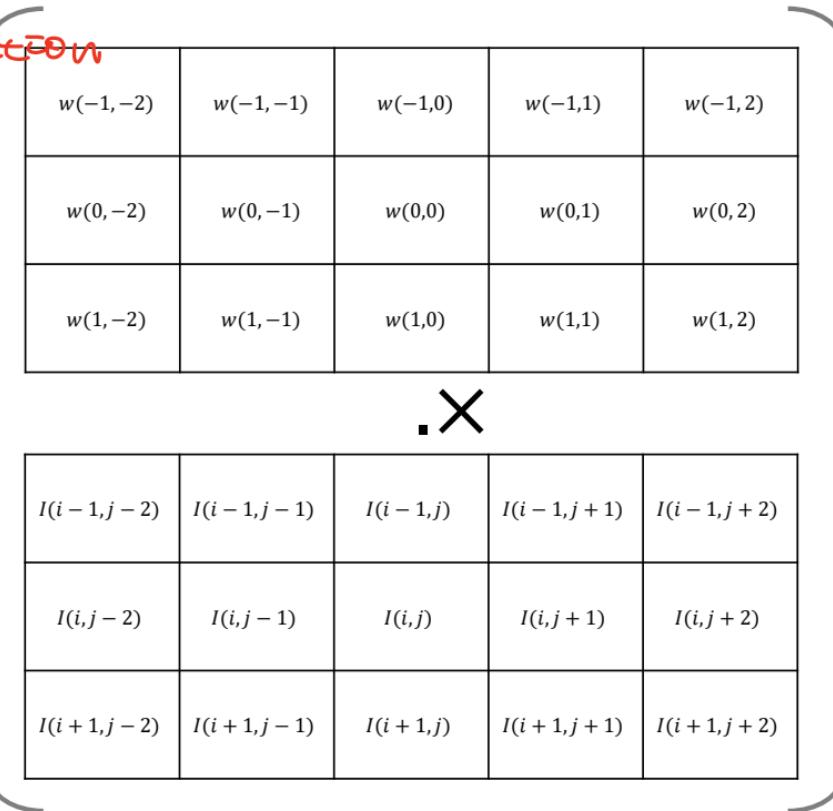


Filtering = SUM (Mask . \times Neighborhood)

(. \times) \downarrow

element wise multiplication

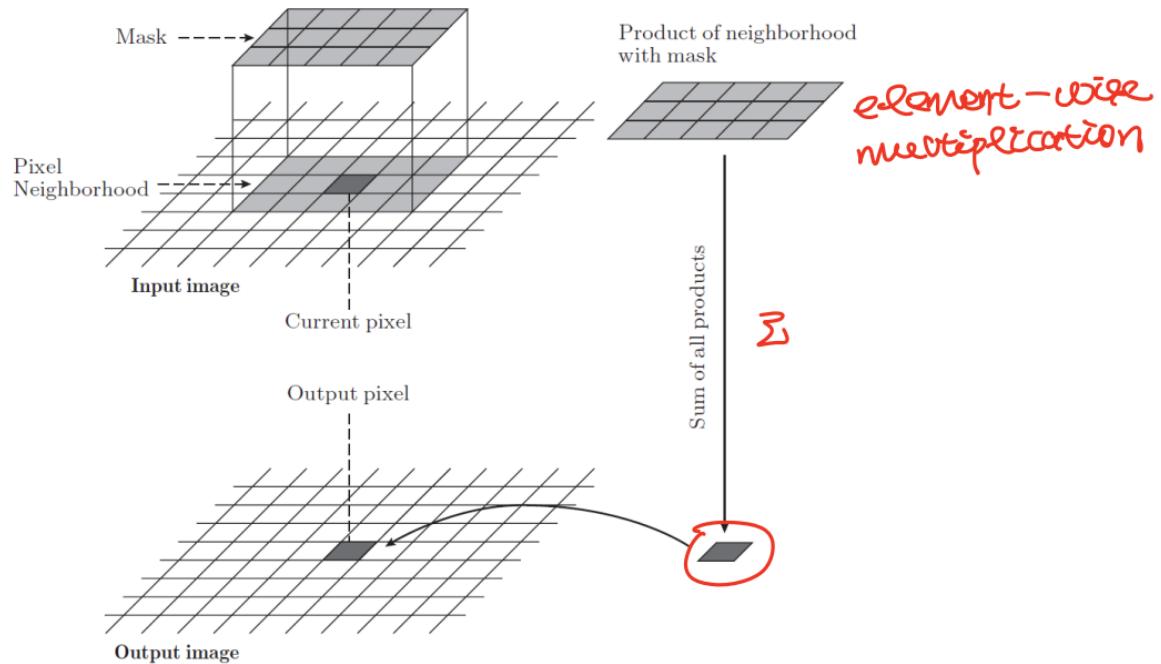
$$O(i, j) = \sum$$



$w(s,t)$

$I(i+s, j+t)$

Filtering = SUM (Mask . \times Neighborhood)



Convolution: From the perspective of signal processing

- **Filtering = Convolution!**
 - When the filter kernel is symmetric
- Let's start with 1-D convolution (from signal processing)
 - Convolution integral

$$y(t) = x(t) * h(t) = \int_{-\infty}^{\infty} x(\tau)h(t - \tau)d\tau \quad \text{For continuous signal}$$

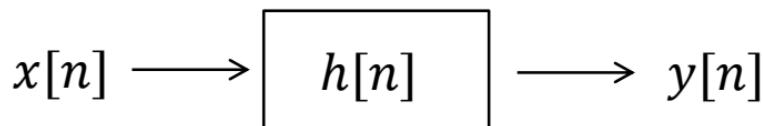
$$y[n] = x[n] * h[n] = \sum_{m=-\infty}^{\infty} x[m]h[n - m] \quad \text{For discrete signal}$$

$$y[n] = \cdots + x[-1]h[n + 1] + x[0]h[n] + x[1]h[n - 1] + \cdots$$

Convolution: From the perspective of signal processing

- **From the perspective of signal processing**
 - An output signal $y[n]$ is obtained by passing an input signal $x[n]$ to a discrete system with a response function $h[n]$.

$$y[n] = x[n] * h[n] = \sum_{m=-\infty}^{\infty} x[m]h[n-m]$$



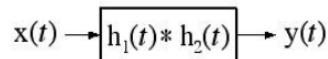
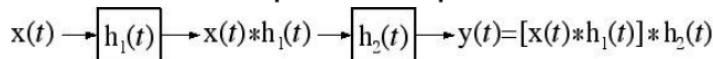
Note **convolution** in spatial domain = **multiplication** in frequency domain

Convolution Properties

- **Associative property**

- Impulse response of a cascade connection
= Convolution of the individual impulse responses

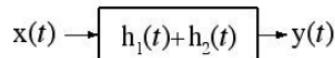
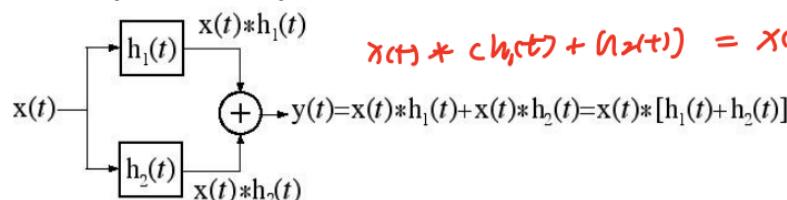
$$h_1(t) * h_2(t) = h_2(t) * h_1(t)$$



- **Distributive property**

- Impulse response of a parallel connection of LTI systems
= Sum of the individual impulse responses.

$$x(t) * (h_1(t) + h_2(t)) = x(t) * h_1(t) + x(t) * h_2(t)$$

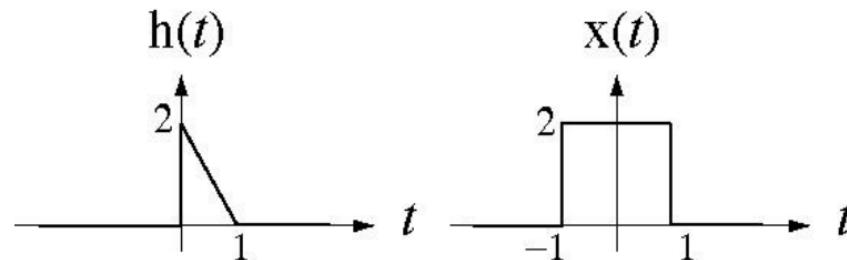


Example

- The convolution on the continuous domain

$$y(t) = x(t) * h(t) = \int_{-\infty}^{\infty} x(\tau)h(t - \tau)d\tau$$

For illustration, let an input signal $x(t)$ and the impulse response $h(t)$ be the two functions below.



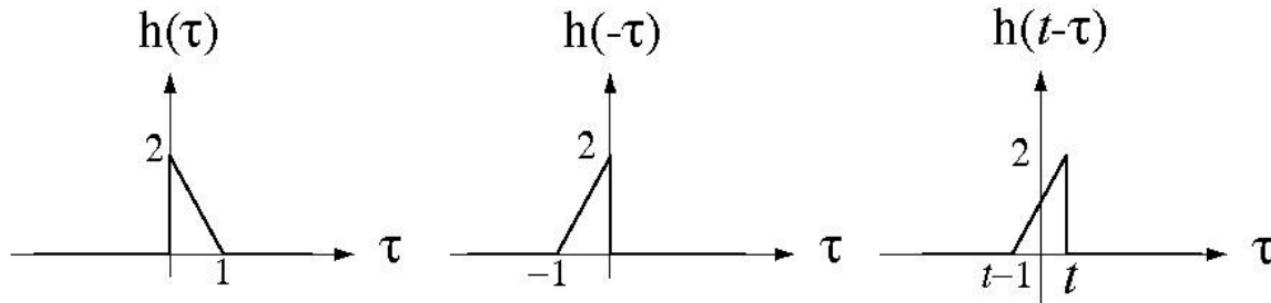
Example

- The convolution on the continuous domain

$$y(t) = x(t) * h(t) = \int_{-\infty}^{\infty} x(\tau)h(t - \tau)d\tau$$

- The functional transformation from $h(t)$ to $h(t - \tau)$

$$h(\tau) \xrightarrow{\tau \rightarrow -\tau} h(-\tau) \xrightarrow{\tau \rightarrow \tau - t} h(-(\tau - t)) = h(t - \tau)$$



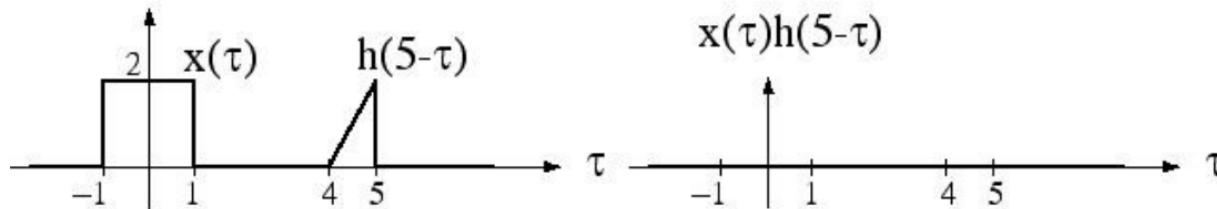
Example

- The convolution on the continuous domain

The convolution value at t :

The area under the product of $x(t)$ and $h(t - \tau)$.

For example, let $t = 5$.



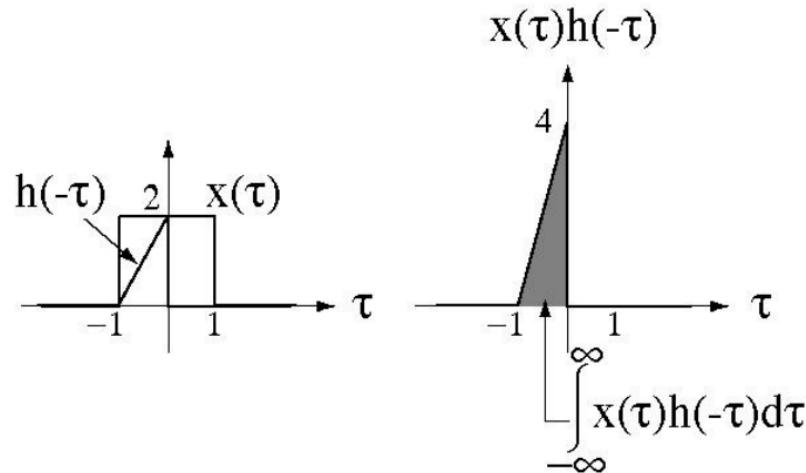
For $t = 5$, the area under the product is zero.

$$\begin{aligned}y(5) &= x(5) * h(5) \\&= \int_{-\infty}^{\infty} x(\tau)h(5 - \tau)d\tau = 0\end{aligned}$$

Example

- The convolution on the continuous domain

When $t = 0$,

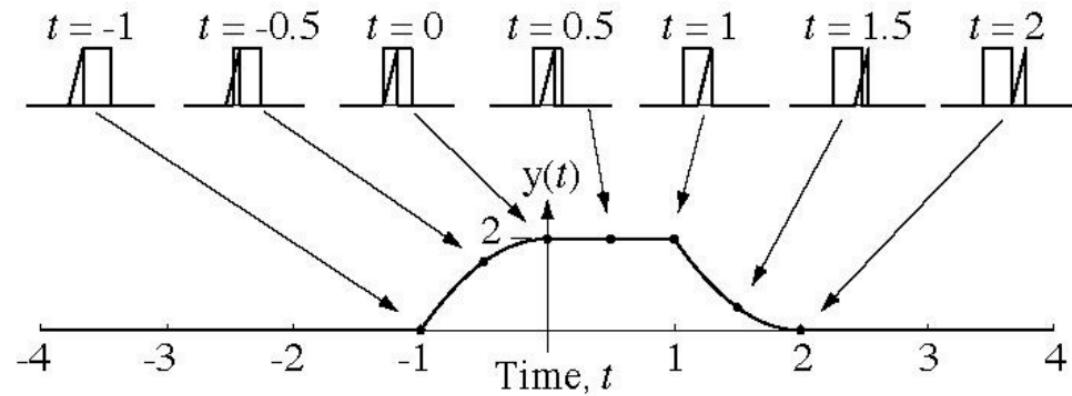


$$y(0) = x(0) * h(0) = \int_{-\infty}^{\infty} x(\tau)h(-\tau)d\tau = 2$$

Example

- The convolution on the continuous domain

The process of convolving to find $y(t)$ is illustrated below.

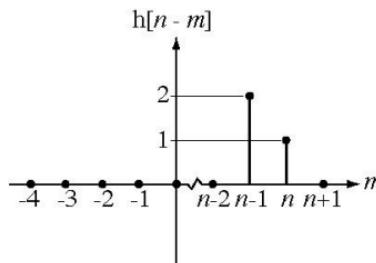
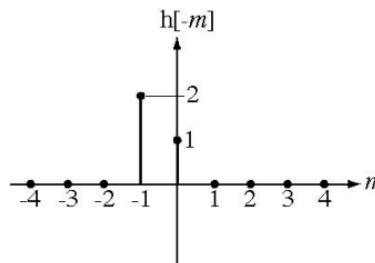
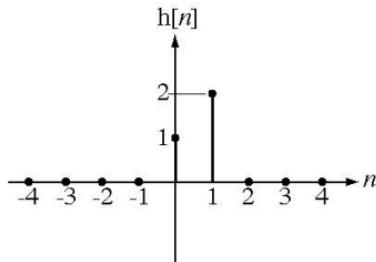
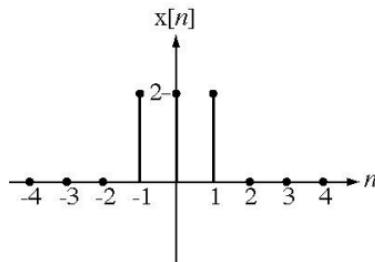


Example

- The convolution on the discrete domain

$$y[n] = x[n] * h[n] = \sum_{m=-\infty}^{\infty} x[m]h[n-m]$$

What is $y[n]$?



Note) $x: M$ points, $h: N$ points
→ $y: M + N - 1$ points

Convolution vs. Filtering

1D convolution

$$\begin{aligned}y[i] &= x[i] * h[i] \\&= \sum_{s=-\infty}^{\infty} h[s]x[i-s]\end{aligned}$$

Note) This formulation is different from previous slides, but the result is identical. Namely, $x * h = h * x$.

1D filtering

$$O[i] = \sum_{s=-a}^a w[s]I[i+s]$$

Remember 2D filtering

$$O(i,j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s,t)I(i+s, j+t)$$

Suppose $h[i]$ exists for $-a \leq i \leq a$ and is symmetric, i.e., $h[-i] = h[i]$. Then, 1D convolution = 1D filtering!
This also applies to N -D convolution and N -D filtering ($N \geq 2$).

2D convolution - example



$$\begin{matrix} * & \begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} & = & ? \end{matrix}$$

2D convolution - example



Original

*

$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix}$$

=



Filtered
(no change)

2D convolution - example

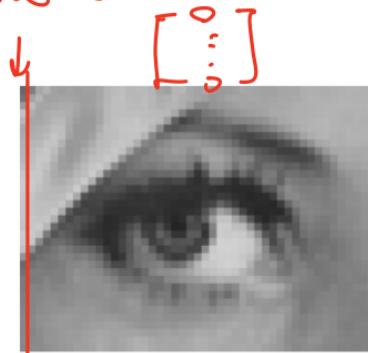


$$\begin{matrix} * & \begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 1 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} & = & ? \end{matrix}$$

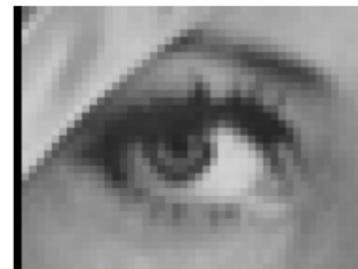
2D convolution - example

In this case, the mask is not symmetrical

+ the column will be all 0 than filtered.



$$\begin{matrix} \textcircled{*} & \uparrow \\ \begin{bmatrix} \cdot 0 & \cdot 0 & \cdot 0 \\ \cdot 0 & \cdot 0 & \cdot 1 \\ \cdot 0 & \cdot 0 & \cdot 0 \end{bmatrix} & = \end{matrix}$$



Shifted right
By 1 pixel

$$\begin{pmatrix} \cdot 0 & \cdot 0 & \cdot 0 \\ \cdot 1 & \cdot 0 & \cdot 0 \\ \cdot 0 & \cdot 0 & \cdot 0 \end{pmatrix}$$

2D convolution - example



$$\text{Original} * \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix} = ?$$

2D convolution - example

average filtering kernel.



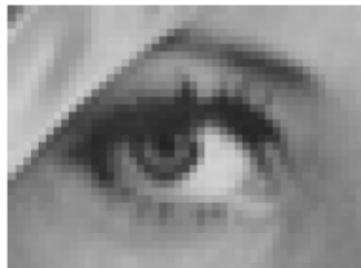
Original

$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} =$$



Blur (with a
box filter)

2D convolution - example



Original

unsharp masking

$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} - \frac{1}{9} \begin{matrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{matrix} = ?$$

(Note that filter sums to 1)

“details of the image”

⇒ *sharpened image*

$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} + \begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} - \frac{1}{9} \begin{matrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{matrix}$$

original - blur = detail

2D convolution - example

- Removing Blurred region means:

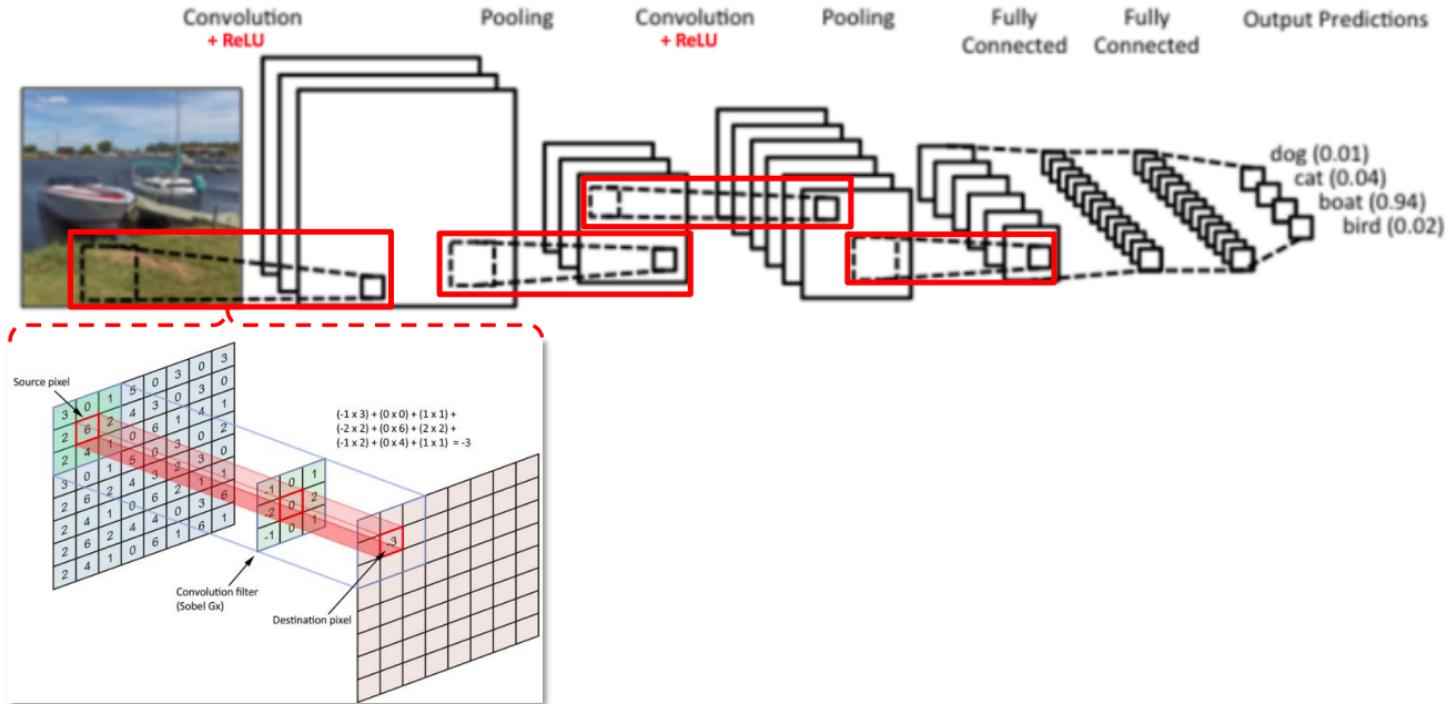


- Adding it back means:



2D convolution - example

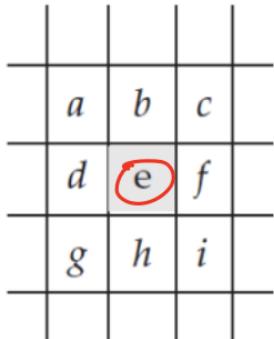
- Deep learning: Convolutional Neural Network



Credit: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

2D Image Filtering: Uniform Mean Filter

- Uniform mean filter
 - The simplest **low-pass** filter



$$\rightarrow \frac{1}{9}(a + b + c + d + e + f + g + h + i)$$

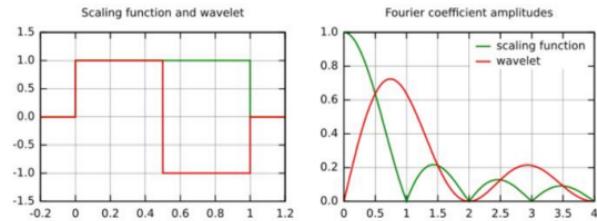
↑
the weights are all the same

Filter kernel

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

EBU6018- remember?

Haar functions in frequency domain



↓
the scaling function
uniform mean filter

Low frequency component passed
high frequency component filtered

2D Image Filtering: Uniform Mean Filter

- How to process pixels at image boundary?
 - 1) Mirroring at an image boundary (mirror padding)

| | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 2 | 2 | 3 | 4 | 5 | 4 | 3 | 2 | 1 | 2 | 3 | 4 | 4 |
| 2 | 2 | 3 | 4 | 5 | 4 | 3 | 2 | 1 | 2 | 3 | 4 | 4 |
| 3 | 3 | | | | | | | | | 3 | 3 | |
| 4 | 4 | | | | | | | | | 4 | 4 | |
| 5 | 5 | | | | | | | | | 5 | 5 | |
| 6 | 6 | | | | | | | | | 6 | 6 | |
| 7 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 2 | 3 | 4 | 5 | 5 |
| 7 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 2 | 3 | 4 | 5 | 5 |

$$O(i, j) = \sum_{s=-1}^1 \sum_{t=-1}^1 w(s, t) I'(i + 1 + s, j + 1 + t)$$

I' : mirror image

2D Image Filtering: Uniform Mean Filter

- How to process pixels at image boundary?
 - 2) Zero padding at an image boundary
 - Common way in convolutional neural networks

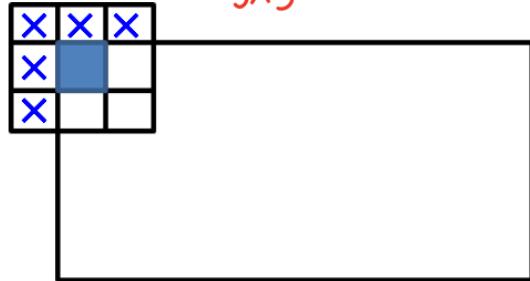
| | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$O(i,j) = \sum_{s=-1}^1 \sum_{t=-1}^1 w(s,t)I'(i + 1 + s, j + 1 + t)$$

I' : zero-padded image

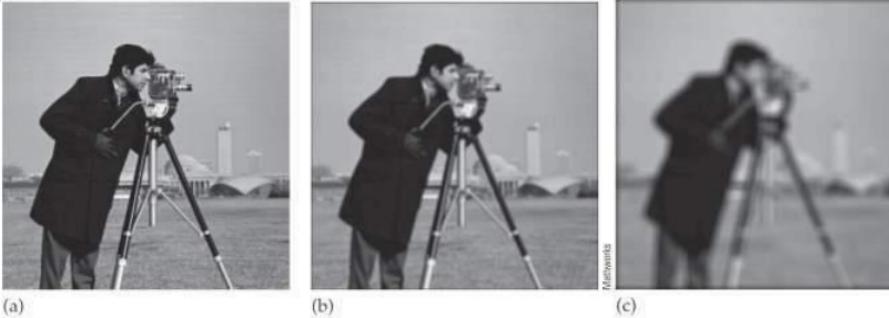
2D Image Filtering: Uniform Mean Filter

- How to process pixels at image boundary?
 - 3) Adjusting filter kernel



$$O(i, j) = \frac{\sum_{s=-1}^1 \sum_{t=-1}^1 w'(s, t) I(i + s, j + t)}{\sum_{s=-1}^1 \sum_{t=-1}^1 w'(s, t)}$$
$$w'(s, t) = (0 \leq i + s \leq H - 1 \& 0 \leq j + t \leq W - 1 ? w(s, t) : 0)$$

2D Image Filtering: Uniform Mean Filter



size $\uparrow \Rightarrow \uparrow$ blurred
(averaging more pixels)

(a) original image, (b) 3×3 filter with zero padding: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$, (c) 9×9 filter with zero padding: $\frac{1}{81} \begin{bmatrix} 1 & \cdots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \cdots & 1 \end{bmatrix}$, (d) 25×25 filter with zero padding: $\frac{1}{625} \begin{bmatrix} 1 & \cdots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \cdots & 1 \end{bmatrix}$, (e) 25×25 filter with mirroring: $\frac{1}{625} \begin{bmatrix} 1 & \cdots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \cdots & 1 \end{bmatrix}$

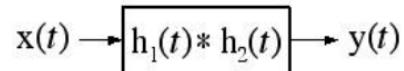
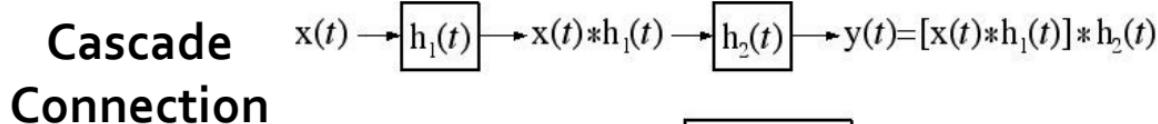
2D Image Filtering: Uniform Mean Filter

- Uniform mean filtering is separable.

- Separable filtering is VERY important in terms of runtime.

↓ computation

$$w(s, t) = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \frac{1}{3} [1 \quad 1 \quad 1] = w_s(s, 0) * w_s(0, t)$$



2D Image Filtering: Gaussian Filter

- Gaussian distribution

- One of the most commonly used parametric models
- Fourier transform (Gaussian func.) = Gaussian func.
- Rotationally symmetric
- Separable filter
- Gaussian * Gaussian = Gaussian

1-D

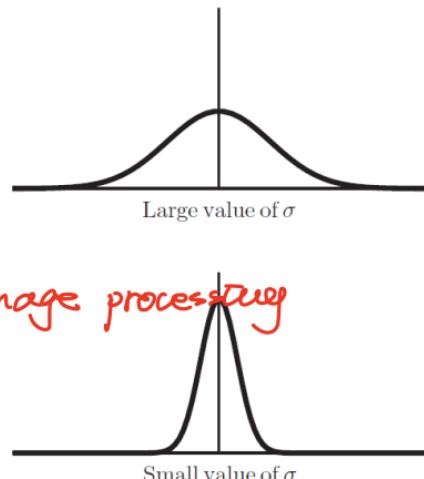
$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

2-D

$$f(x, y) = \frac{1}{2\pi\sigma_x\sigma_y} \exp\left(-\frac{(x - \mu_x)^2}{2\sigma_x^2} - \frac{(y - \mu_y)^2}{2\sigma_y^2}\right)$$

N-D

$$f(\mathbf{m}) = \frac{1}{(2\pi)^{N/2} |\Sigma|^{1/2}} \exp(-(\mathbf{m} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{m} - \boldsymbol{\mu}))$$



Q: What if σ approaches infinite?

FIGURE 5.1 One-dimensional Gaussians.

2D Image Filtering: Gaussian Filter

- Gaussian filter's advantage

- It considers spatial distances within neighborhoods
- Blurred results looks more natural compared to mean filter.



5 × 5, $\sigma = 0.5$



5 × 5, $\sigma = 2$



11 × 11, $\sigma = 1$



11 × 11, $\sigma = 5$

Practical implementation

$$O(i, j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) I(i + s, j + t)$$

Zero mean Gaussian filter
 $w(s, t)$

$$= \frac{1}{\sum_{m=-a}^a \sum_{n=-b}^b \exp\left(-\frac{m^2}{2\sigma_s^2} - \frac{n^2}{2\sigma_t^2}\right)} \exp\left(-\frac{s^2}{2\sigma_s^2} - \frac{t^2}{2\sigma_t^2}\right)$$

In results, $\mu_s = \mu_t = 0$ and $\sigma_s = \sigma_t = \sigma$

2D Image Filtering: Gaussian Filter

- Gaussian filter's advantage

- Separable

Practical implementation

$$O(i, j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) I(i + s, j + t)$$



$$O(i, j) = \sum_{s=-a}^a w_s(s) \sum_{t=-b}^b w_t(t) I(i + s, j + t)$$

$$\begin{aligned} & \frac{1}{2\pi\sigma_s\sigma_t} \exp\left(-\frac{s^2}{2\sigma_s^2} - \frac{t^2}{2\sigma_t^2}\right) \\ &= \frac{1}{\sqrt{2\pi}\sigma_s} \exp\left(-\frac{s^2}{2\sigma_s^2}\right) \frac{1}{\sqrt{2\pi}\sigma_t} \exp\left(-\frac{t^2}{2\sigma_t^2}\right) \\ & w(s, t) \\ &= \frac{1}{\sum_{m=-a}^a \sum_{n=-b}^b \exp\left(-\frac{m^2}{2\sigma_s^2} - \frac{n^2}{2\sigma_t^2}\right)} \exp\left(-\frac{s^2}{2\sigma_s^2} - \frac{t^2}{2\sigma_t^2}\right) \\ & w_s(s) = \frac{1}{\sum_{m=-a}^a \exp\left(-\frac{m^2}{2\sigma_s^2}\right)} \exp\left(-\frac{s^2}{2\sigma_s^2}\right) \\ & w_t(t) = \frac{1}{\sum_{n=-b}^b \exp\left(-\frac{n^2}{2\sigma_t^2}\right)} \exp\left(-\frac{t^2}{2\sigma_t^2}\right) \end{aligned}$$

Frequency: Low- and High-pass Filters

- Frequency in an image
 - How often do intensity values vary in neighbourhoods?
- Low-pass filter: uniform average filter, Gaussian filter
- High-pass filter: Sobel filter, Laplacian filter
 - (this estimates intensity change.)

Example of Low-pass filter: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

Example of High-pass filter: $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$

2D Image Filtering: Sobel Filter

- Using the first order gradient

$$\nabla I = (I_x, I_y) = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

Practical implementation

$$\nabla I = (I_x, I_y) = (|S_x * I|, |S_y * I|)$$

For color image,

$$M(x, y) = (M_R(x, y) + M_G(x, y) + M_B(x, y))/3$$

Sobel filter output $M(x, y) = \sqrt{|I_x|^2 + |I_y|^2}$ or $|I_x| + |I_y|$

$$\begin{matrix} S_x & & S_y \\ \begin{array}{|c|c|c|} \hline -1 & -2 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 2 & 1 \\ \hline \end{array} & & \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} \end{matrix}$$

| | | |
|----|----|----|
| -1 | -2 | -1 |
| 0 | 0 | 0 |
| 1 | 2 | 1 |

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

2D Image Filtering: Laplacian Filter

- Using the second order gradient

- 1st derivative of an image $I(x, y)$

$$\nabla I(x, y) = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

$$I_x(x, y) = \frac{\partial I}{\partial x} = I(x + 1, y) - I(x, y) \quad I_y(x, y) = \frac{\partial I}{\partial y} = I(x, y + 1) - I(x, y)$$

- 2nd derivative of an image $I(x, y)$

$$\nabla^2 I(x, y) = \left(\frac{\partial^2 I}{\partial x^2}, \frac{\partial^2 I}{\partial y^2} \right)$$

$$\frac{\partial^2 I}{\partial x^2} = I_x(x, y) - I_x(x - 1, y) = I(x + 1, y) + I(x - 1, y) - 2I(x, y)$$

$$\frac{\partial^2 I}{\partial y^2} = I_y(x, y) - I_y(x, y - 1) = I(x, y + 1) + I(x, y - 1) - 2I(x, y)$$

$$\text{Laplacian Filtering output } G(x, y) = \left| \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2} \right|$$

2D Image Filtering: Laplacian Filter

- Using the second order gradient

Laplacian filter L

| | | |
|---|----|---|
| 0 | 1 | 0 |
| 1 | -4 | 1 |
| 0 | 1 | 0 |

or

| | | |
|---|----|---|
| 1 | 1 | 1 |
| 1 | -8 | 1 |
| 1 | 1 | 1 |

$$O = |L * I|$$

More general form

* the sum of
the weights in
Laplacian mask
should be $\underline{\underline{0}}$

| | | |
|-----------------------------|-----------------------------|-----------------------------|
| $\frac{\alpha}{1+\alpha}$ | $\frac{1-\alpha}{1+\alpha}$ | $\frac{\alpha}{1+\alpha}$ |
| $\frac{1-\alpha}{1+\alpha}$ | $\frac{-4}{1+\alpha}$ | $\frac{1-\alpha}{1+\alpha}$ |
| $\frac{\alpha}{1+\alpha}$ | $\frac{1-\alpha}{1+\alpha}$ | $\frac{\alpha}{1+\alpha}$ |



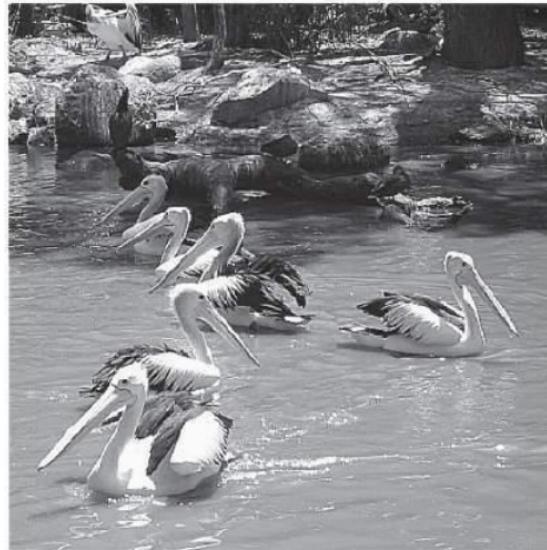
For color image,
 $O(x, y) = (O_R(x, y) + O_G(x, y) + O_B(x, y))/3$

Unsharp Masking

- Makes an image look sharper by boosting high-frequency components.

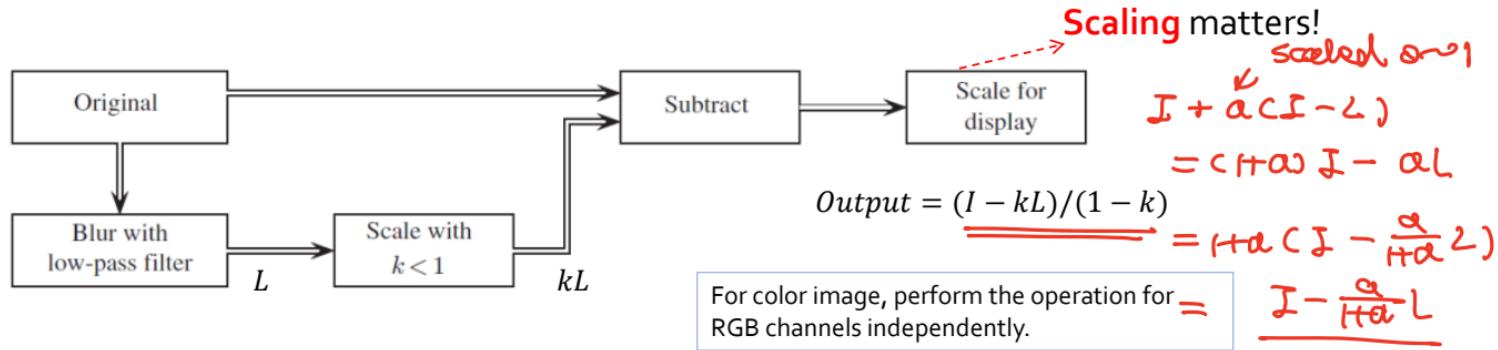


Original
image



Unsharp mask
result

Unsharp Masking



```
x = double(imread('cameraman.tif'))/255;  
  
f = fspecial('average');  
xf = filter2(f,x);  
figure, imshow(xf)  
  
% k: parameter determining the amount of reducing low-frequency component  
% 0<=k<=1  
% As k increases, the output image looks sharper. But, it is recommended using k below 0.5;  
k = 0.5;  
fi = zeros(3); fi(2,2)=1;  
f2 = (fi - k*f)/(1-k);  
xf2 = filter2(f2,x);  
figure, imshow(xf2)
```

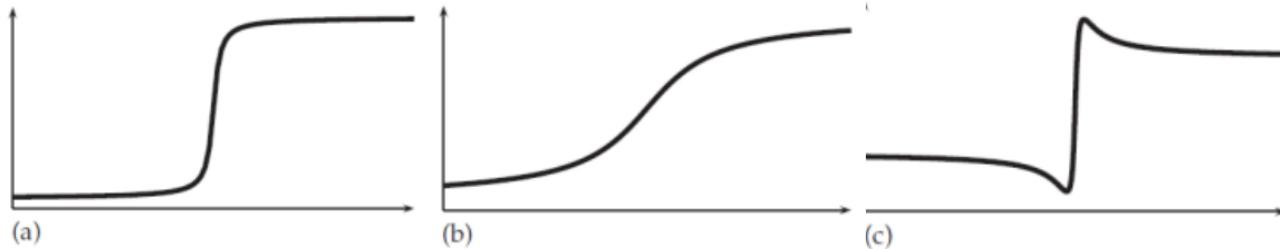
Matlab Code

$$= \frac{I - FL}{1 - k}$$

$$k = \frac{\alpha}{1 + \alpha}$$

$$\alpha \in (0,1) \Rightarrow k \in (0, \frac{1}{2})$$

Unsharp Masking



$$Output = (I - kL)/(1 - k)$$

$$w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - k \begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

↓ **Scaling**

$$w = \frac{1}{1-k} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{k}{1-k} \begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

2D Image Filtering: Nonlinear Filter

- **Linear vs. Nonlinear filters?**
 - Definition of linearity, and its advantages
- **Some instances of nonlinear filter**
 - Max, Min, Median filter



Original



Max filter



Min filter



Median filter

2D Image Filtering: Nonlinear Filter

- **Median Filter**

1. Sort pixels within the window centered at reference pixel
2. Select the median value as an output.

Q: Mean vs. Median?



Original



Max filter



Min filter



Median filter

Summary

- **Spatial Filtering**
 - Depends on how to define a filter kernel.
 - Linear vs. Nonlinear filter: Linear filter is **separable!** *(+ computation)*
 - Mean, *non-linear* (median, maximum, minimum), low-pass, high-pass, Gaussian, Laplacian filter
 - Edge Sharpening: combination of low-pass and high-pass filters.
- **Convolution**
 - Same as filter *(symmetrical)*
 - Commutativity, Cascade property, Parallel property

EBU7240

Computer Vision

- Restoration: noise and noise removal -

Semester 1, 2021

Changjae Oh

Contents

- **Image Degradation Model**
 - Image Noise
- **Noise removal**
 - Salt-and-Pepper Noise Removal
 - Gaussian Noise Removal
 - Periodic Noise Removal

Image Degradation Model

- **Image restoration**
 - Aims to reduce the image degradation
- **Types of image degradation**
 - Noise, out-of-focus blur, motion blur

low - light
underwater
haze



Image Degradation Model

- When an input degraded image $g(x, y)$ is given, our goal is to estimate $f(x, y)$.

$$g(x, y) = h(x, y) * f(x, y) + n(x, y)$$

clean image

- $n(x, y)$: Additive noise
- $h(x, y)$: Blur kernel, which is the same as filtering mask.
- In frequency domain (using 2D DFT)?

$$G(u, v) = H(u, v)F(u, v) + N(u, v)$$



$$F(u, v) = (G(u, v) - \underline{N(u, v)}) / \underline{H(u, v)}$$

(x, y): 2D image coordinate
(u, v): 2D frequency coord.

Is that it? It is not such a simple problem.

- 1) $N(u, v)$ is not known.
- 2) What if $H(u, v) = 0$?

Image Noise

- **Noise**
 - Any kind of degradation in an image caused by external disturbance
 - To make the problem simple, we pre-assume the noise models.
The most appropriate restoration may vary depending on the noise types.
- **Noise types**
 - Periodic noise, Salt and Pepper noise, Gaussian noise, Speckle noise
- In the beginning, let's think of the case using $h(x, y)=1$

$$g(x, y) = f(x, y) + n(x, y)$$

↓
Don't have blur

Image Noise: Salt and Pepper Noise

- Sharp and sudden disturbances
- Image is **randomly** scattered as white (salt) or black (pepper) pixels.



(a)



(b)

FIGURE 8.1 Noise on an image. (a) Original image. (b) With added salt and pepper noise.

Image Noise: Gaussian Noise

- Additive White Gaussian Noise (AWGN)

- **Additive** noise: $I_G(x, y) = I(x, y) + N(x, y)$
 - **White** noise: randomly fluctuated and normally distributed
 - Most approaches assume this type of noise.

- Usually, zero mean AWGN is assumed ($\mu = 0$).

The probability density function (PDF) P of a Gaussian random variable z is

$$P(z = N) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(z - \mu)^2}{2\sigma^2}\right)$$



the distribution of the noise value is \Rightarrow Gaussian distribution



Image Noise: Speckle Noise

- **Multiplicative Noise**

- $I_1(x, y) = I(x, y) + I(x, y)N(x, y)$
- $N(x, y)$: zero mean uniform distributed function with σ
- This noise is usually in the active radar, synthetic aperture radar (SAR), medical ultrasound and optical coherence tomography images.
- The reduction of Speckle noise is MUCH more difficult due to the multiplication

Image Noise: Gaussian Noise vs. Speckle Noise



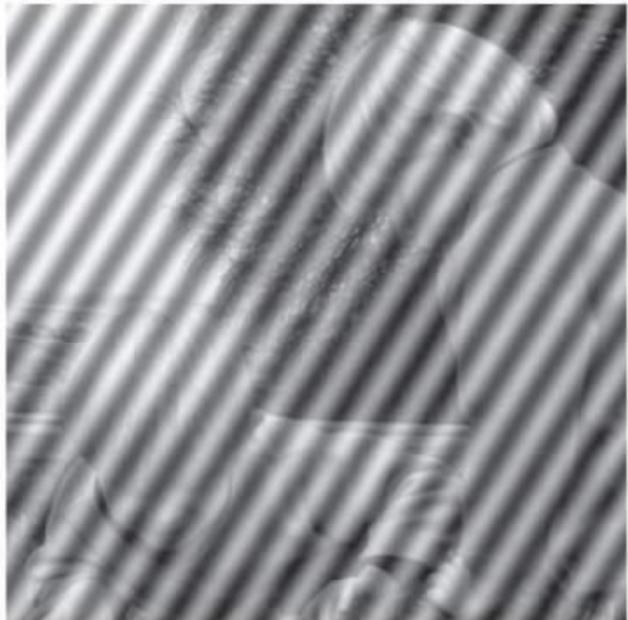
Gaussian Noise



Speckle Noise

Image Noise: Periodic Noise

- **Periodic fluctuation**
 - Electrical or electromechanical interference during image acquisition
 - Spatially dependent noise
 - It can be modeled as sinusoid waves



Contents

- Image Degradation Model
 - Image Noise
- **Noise removal**
 - Salt-and-Pepper Noise Removal
 - Gaussian Noise Removal
 - Periodic Noise Removal

Salt and Pepper Noise Removal

- Low-Pass Filtering
 - For instance, uniform averaging filter or Gaussian averaging filter
 - Not so effective
- Median Filtering
 - Median filter works well in the salt-and-pepper noise removal.
 - Rank-order Filtering is a general formulation of median filter.
Sort the pixel intensity everytime
↓
long time process
- Outlier Rejection Method
(sampler method)
↓ computation

Salt and Pepper Noise Removal

- Low-Pass Filtering

- NOT effective in removing the salt and pepper noise.

filter size ↑ → ↑ noise remove
Blur the noise



Original Image



10% salt and pepper noise
(salt: 5%, pepper: 5%)



3x3 average filter

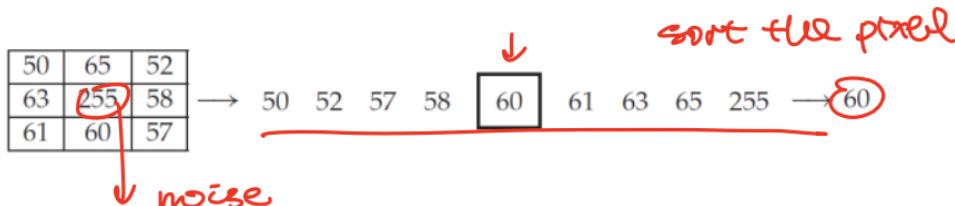


7x7 average filter

Salt and Pepper Noise Removal: Median Filtering

- Median filtering procedure

- Sort pixels within the window centered at reference pixel
- Select the median value as an output.



- Good performance in the salt-and-pepper noise removal

Why?

$$\begin{matrix} 50 & 65 & 52 \\ 63 & \underline{60} & 58 \\ 61 & 60 & 57 \end{matrix}$$



3 × 3 median filter

Salt and Pepper Noise Removal: Median Filtering

- More results using median filter



20% salt and pepper noise
(salt: 10%, pepper: 10%)



Result using 3×3 median filter



Result applying 3×3 median filter twice



Result applying 5×5 median filter

Salt and Pepper Noise Removal: Outlier Rejection Method

- The brute force implementation of median filtering is very slow.
 - Different method was proposed to remove the salt-and-pepper noise efficiently.

- Key idea: **Outlier detection → rejection**

- Outliers usually tend to be different from neighboring pixels' intensities.

↪ observation



- Outlier rejection method

- Choose a threshold value D



$$m = 0 \\ p = 2 \rightarrow \text{noisy.} \Rightarrow$$



- For a given pixel, compare its value p with the mean m of the values of its eight neighbors

- If $|p - m| > D$, then classify the pixel as noisy, otherwise not.

- If the pixel is noisy, replace its value with m . ← ? not median

Salt and Pepper Noise Removal: Outlier Rejection Method



(a)



(b)

FIGURE 8.9 Applying the outlier method to 10% salt and pepper noise. (a) $D = 0.2$.
(b) $D = 0.4$.

Gaussian Noise Removal

- Why does an image average filter work well for Gaussian Noise removal?
- Additive White Gaussian Noise (AWGN) $N(x, y)$ is added as below.
 - $I_G(x, y) = I(x, y) + N(x, y)$

1. Suppose we have 100 noisy images I_G .

- $I_G^i(x, y) = I(x, y) + N^i(x, y) \quad i = 1, \dots, 100$

2. Average 100 noisy images

$$\frac{1}{100} \sum_{i=1}^{100} I_G^i(x, y) = I(x, y) + \frac{1}{100} \sum_{i=1}^{100} N^i(x, y)$$

can be also considered as
10x10 kernel Gaussian
10x10 kernel filtering

//

This will be close to 0, as AWGN N is a zero-mean Gaussian PDF.

Gaussian Noise Removal

- Why does an image average filter work well for Gaussian Noise removal?

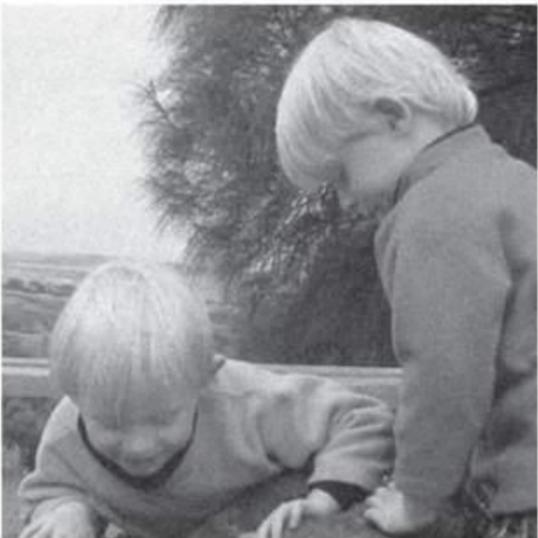


Image averaging to remove Gaussian noise. (a) 10 images (b) 100 images

Gaussian Noise Removal: Simple Average Filtering

- **Simple Average Filtering: Uniform mean filter or Gaussian filter**
 - Using a small window: not so effective in noise removal.
 - Using a large window: effective in noise removal, but the output is over-smoothed.



Uniform mean filter to remove Gaussian noise. (a) 3×3 averaging (b) 5×5 averaging

Gaussian Noise Removal: Bilateral Filtering

- Bilateral filter for grayscale image

- One of the most popular filters with various applications
- Considers both spatial and intensity distances

$$O(i,j) = \sum_{s=-a}^a \sum_{t=-b}^b w(s,t) I(i+s, j+t)$$

spatial term *Intensity term*

$$w(s,t) = \frac{1}{W(i,j)} \exp\left(-\frac{s^2}{2\sigma_s^2} - \frac{t^2}{2\sigma_t^2}\right) \exp\left(-\frac{(I(i,j) - I(i+s, j+t))^2}{2\sigma_r^2}\right)$$
$$W(i,j) = \sum_{m=-a}^a \sum_{n=-b}^b \exp\left(-\frac{m^2}{2\sigma_s^2} - \frac{n^2}{2\sigma_t^2}\right) \exp\left(-\frac{(I(i,j) - I(i+m, j+n))^2}{2\sigma_r^2}\right)$$

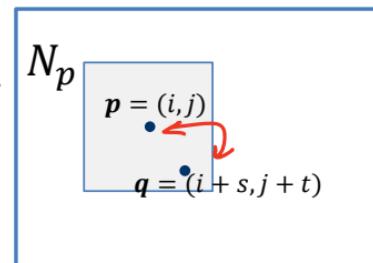
Gaussian

- This can be rewritten as:

$$O_p = \frac{1}{W_p} \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|I_p - I_q|) I_q$$

$$W_p = \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|I_p - I_q|)$$

Gaussian filter: spatial distance
Bilateral filter: intensity distance



{ pixel intensity similar \Rightarrow high w
{ pixel intensity diff \Rightarrow low w

Gaussian Noise Removal: Bilateral Filtering

- Bilateral filter for color image

- Applying filter to each channel

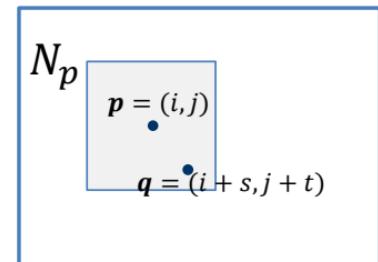
$$C_p = (R_p, G_p, B_p)^T$$

$$R_p = \frac{1}{W_p} \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|C_p - C_q|) R_q$$

$$G_p = \frac{1}{W_p} \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|C_p - C_q|) G_q$$

$$B_p = \frac{1}{W_p} \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|C_p - C_q|) B_q$$

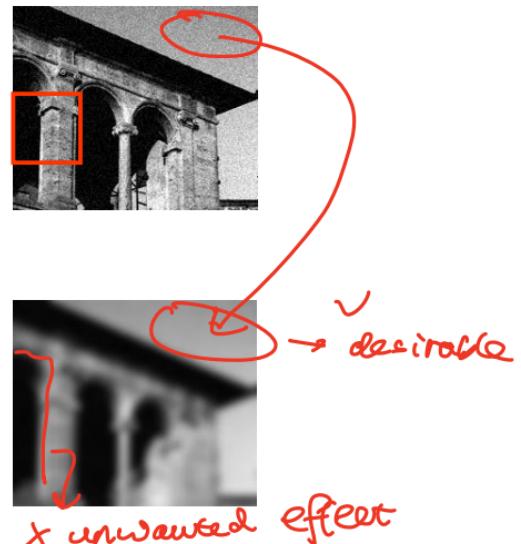
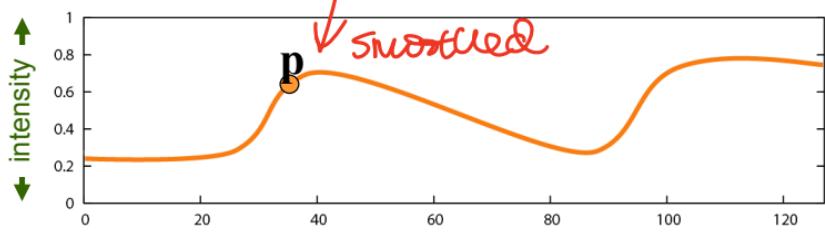
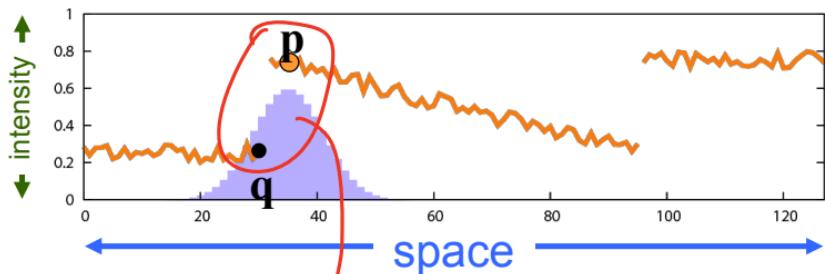
$$W_p = \sum_{q \in N_p} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|C_p - C_q|)$$



Gaussian filter vs. Bilateral filter

- Gaussian filter

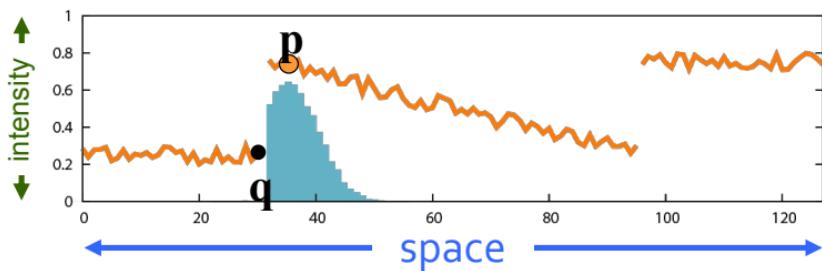
- Weighted average of neighbors
- Depends only on spatial distance
- No edge term



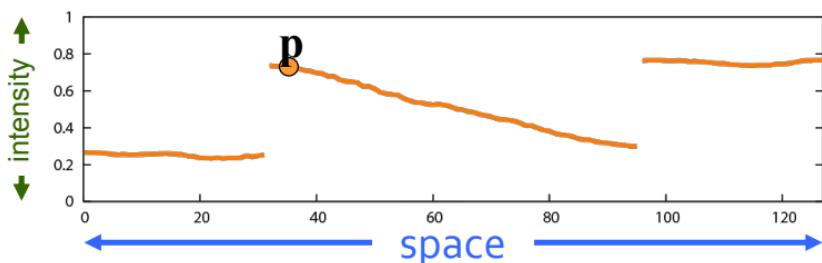
Gaussian filter vs. Bilateral filter

- **Bilateral filter**

- Weighted average of neighbors
- Depends on spatial and range difference

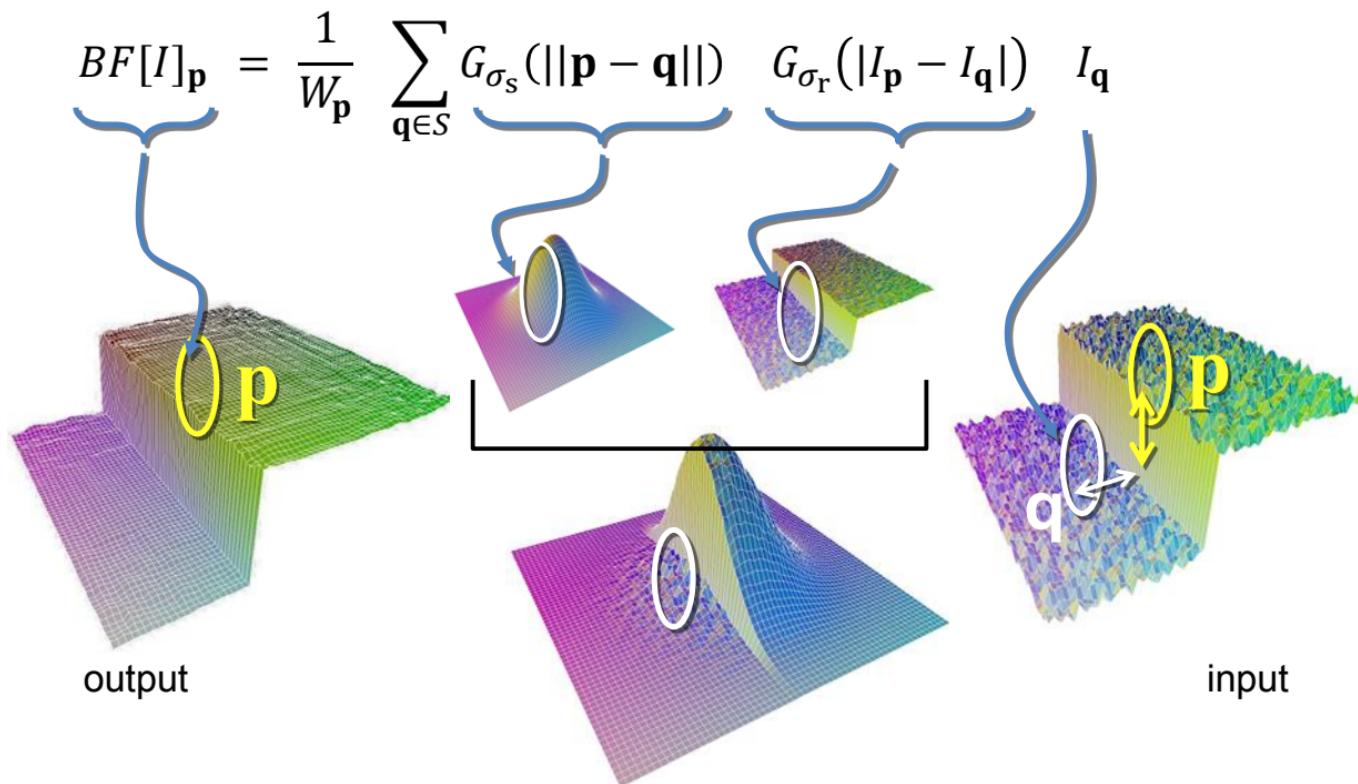


Input



Output

Bilateral filter on a height field



reproduced
from [Durand 02]

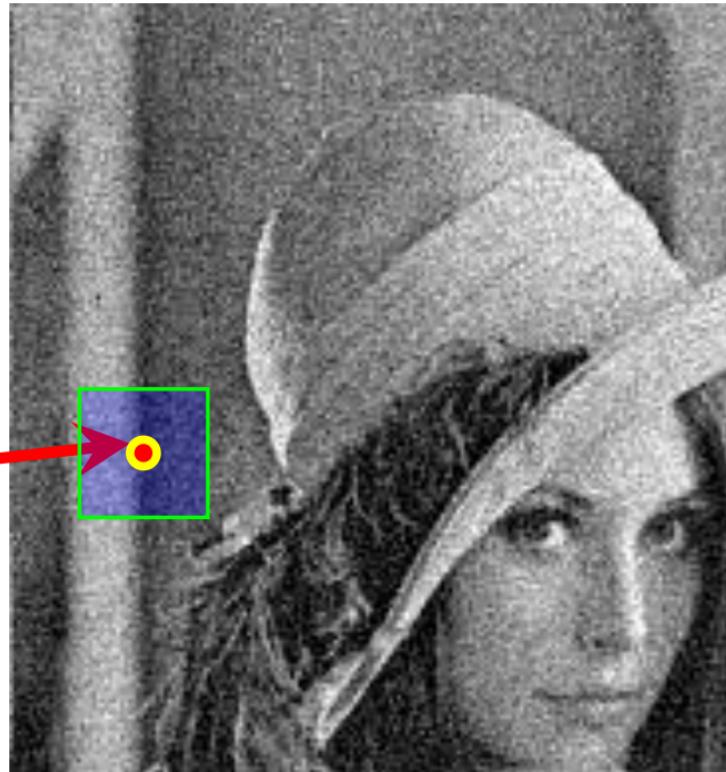
Gaussian Noise Removal: Non-local Means Filtering

- **Same goals:**
 - Smooth within Similar Regions
- **KEY INSIGHT:**
 - Generalize, extend 'Similarity'
- **Bilateral:**
 - Averages neighbors with similar intensities;
- **NL-Means:**
 - Averages neighbors with similar neighborhoods!

Gaussian Noise Removal: Non-local Means Filtering

- For each pixel p :
 - Define a small, simple fixed size neighborhood;
 - Define vector V_p : a list of neighboring pixel values.

$$V_p = \begin{bmatrix} 0.74 \\ 0.32 \\ 0.41 \\ 0.55 \\ \dots \\ \dots \\ \dots \end{bmatrix}$$

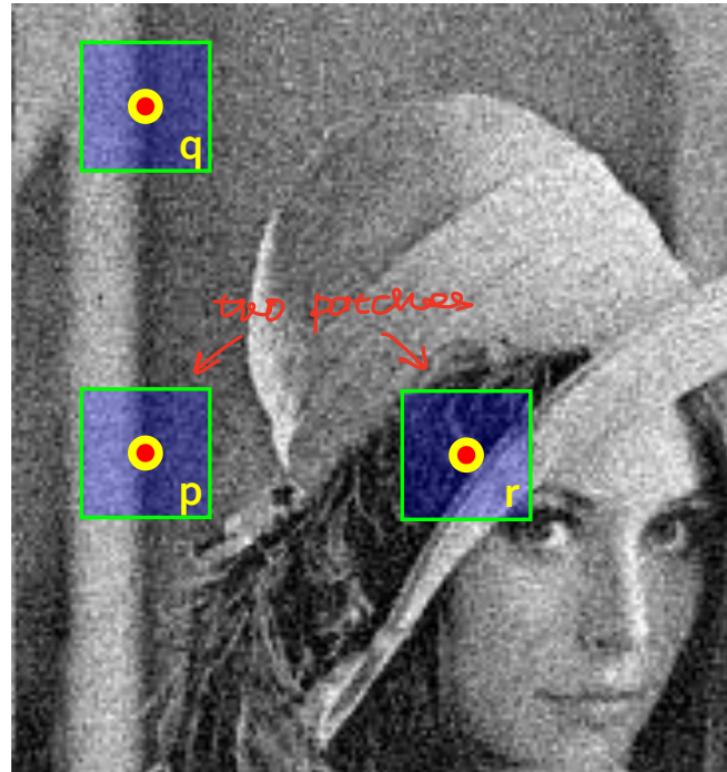


Gaussian Noise Removal: Non-local Means Filtering

- For each pixel p :

- Define a small, simple fixed size neighborhood;
- Define vector V_p : a list of neighboring pixel values.
- 'Similar' pixels $p, q \rightarrow$ SMALL distance $\|V_p - V_q\|_2$
- 'Dissimilar' pixels $p, r \rightarrow$ LARGE distance $\|V_p - V_r\|_2$

over the whole image
(non-local)



Gaussian Noise Removal: Non-local Means Filtering

- For each pixel p :
 - Define a small, simple fixed size neighborhood;
 - Define vector \mathbf{V}_p : a list of neighboring pixel values.
 - 'Similar' pixels $p, q \rightarrow$ SMALL distance $\|\mathbf{V}_p - \mathbf{V}_q\|_2$
 - 'Dissimilar' pixels $p, r \rightarrow$ LARGE distance $\|\mathbf{V}_p - \mathbf{V}_r\|_2$
- Filtering with this neighboring pixels!
 - No spatial terms,
 - Measures the distance between patches (neighbor pixels)

$$NLMF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(\|\vec{V}_p - \vec{V}_q\|^2) I_q$$



Gaussian Noise Removal: Non-local Means Filtering

- Noisy source image:



Gaussian Noise Removal: Non-local Means Filtering

- Gaussian Filter
 - Low noise
 - Low detail



Gaussian Noise Removal: Non-local Means Filtering

- **Bilateral Filter**

- Better at noise removal
 - but 'stairsteps'



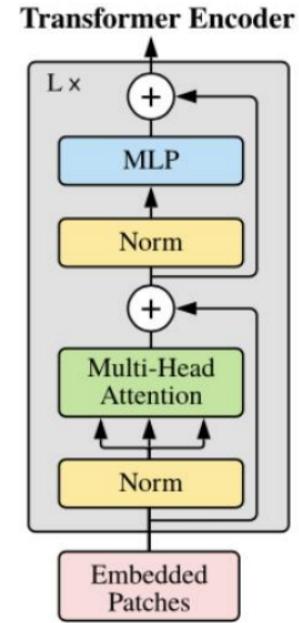
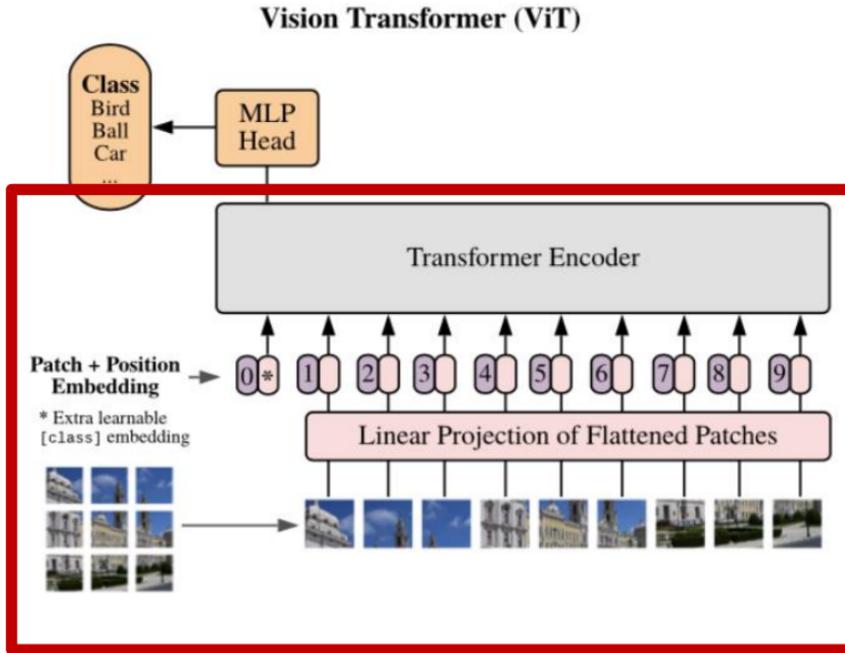
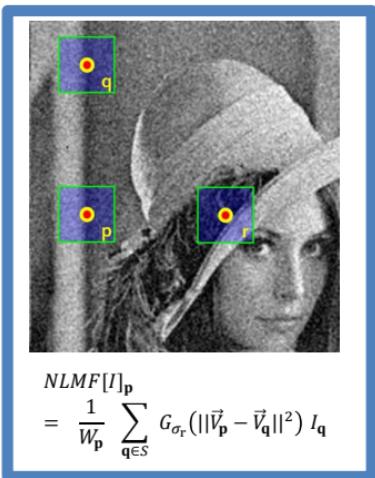
Gaussian Noise Removal: Non-local Means Filtering

- **NL-Means**
 - Sharp
 - Low noise
 - Few artifacts



Non-local Means Filtering: Old-fashioned?

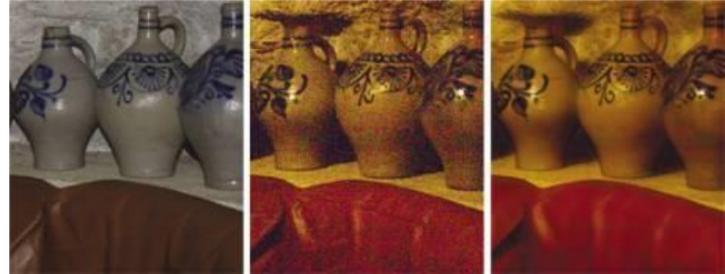
- Vision Transformer (in deep learning)
 - A new architecture that shows state-of-the-art performance in computer vision tasks



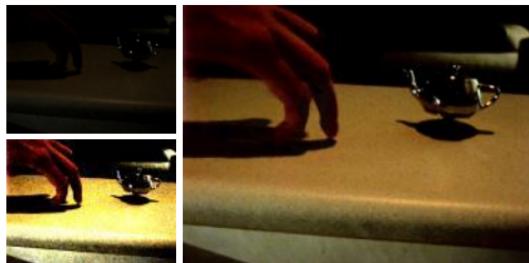
Many applications, not limited to denoising



Tone Mapping [Durand 02]



Flash / No-Flash [Eisemann 04, Petschnigg 04]

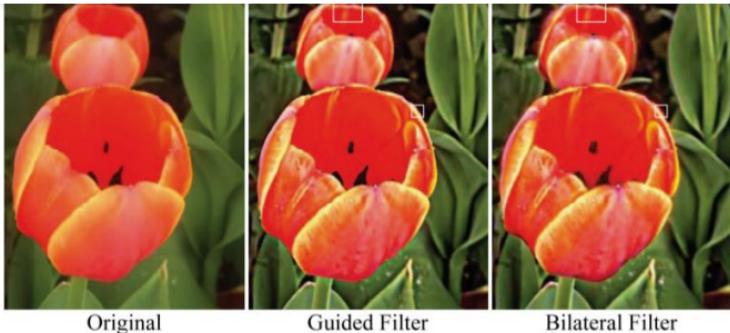


Virtual Video Exposure [Bennett 05]

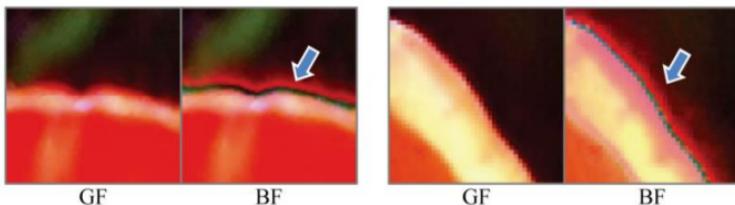


Tone Management [Bae 06]

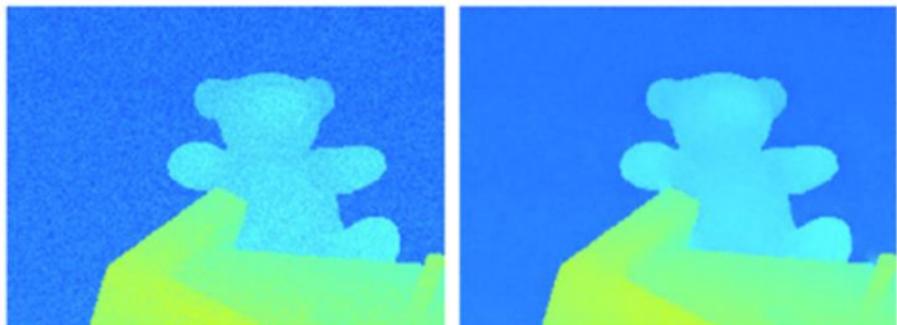
Edge-aware smoothing filters



[Li 2011]



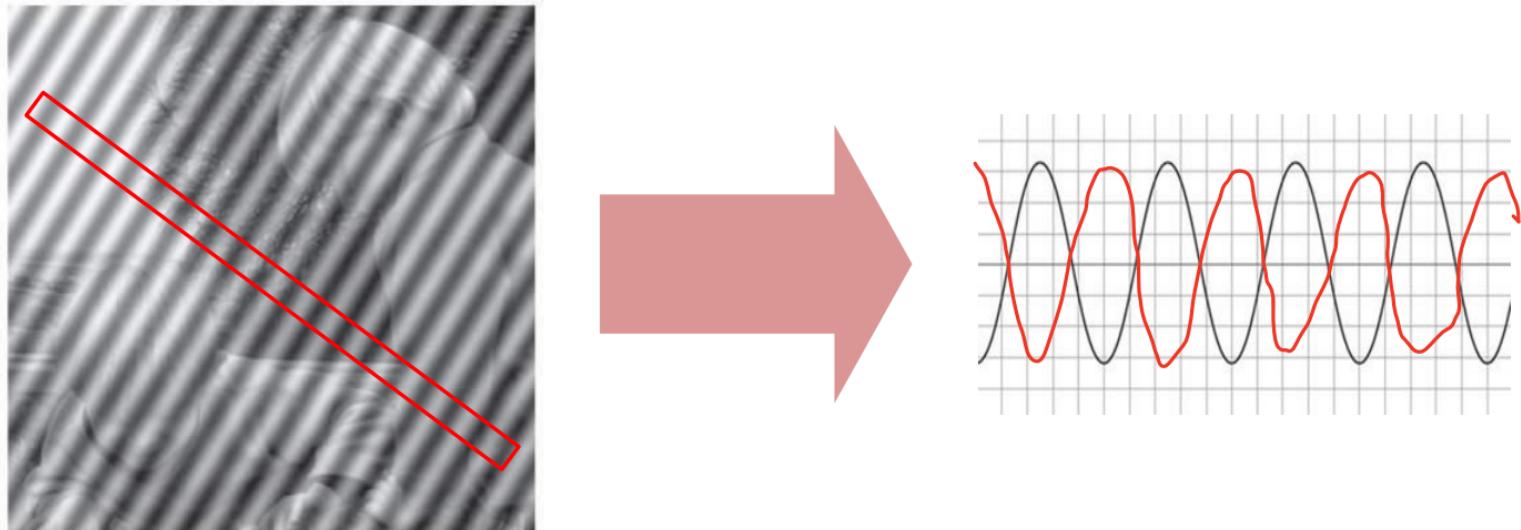
[He 2010]



[Shen 2015]

Periodic Noise Removal: Frequency Domain Filtering

- Periodic fluctuation can be modeled as sinusoid waves



Periodic Noise Removal: Frequency Domain Filtering

- Periodic fluctuation can be modeled as sinusoid waves

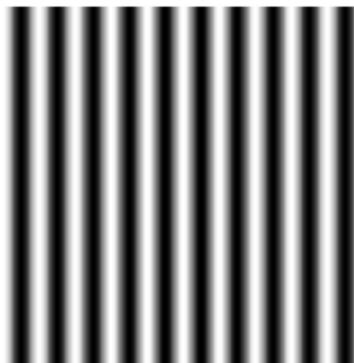
- Spatial domain

$$r(x, y) = A \sin \left[\frac{2\pi u_0(x + B_x)}{M} + \frac{2\pi v_0(y + B_y)}{N} \right]$$

A : amplitude

u_0, v_0 : sinusoidal frequencies

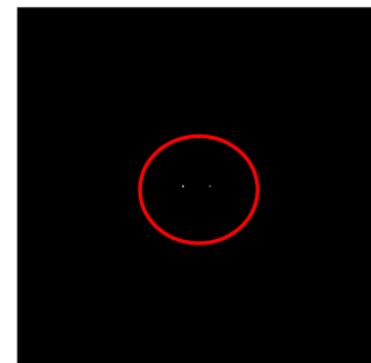
B_x, B_y : phase displacement w.r.t. the origin



Sinusoidal wave in
the spatial domain

- Frequency domain

$$\begin{aligned} R(u, v) &= j \frac{AMN}{2} \left[e^{-j2\pi\left(\frac{u_0B_x}{M} + \frac{v_0B_y}{N}\right)} \delta(u + u_0, v + v_0) \right. \\ &\quad \left. - e^{j2\pi\left(\frac{u_0B_x}{M} + \frac{v_0B_y}{N}\right)} \delta(u - u_0, v - v_0) \right] \end{aligned}$$



Two symmetric spikes
in the Fourier domain

Periodic Noise Removal: Frequency Domain Filtering

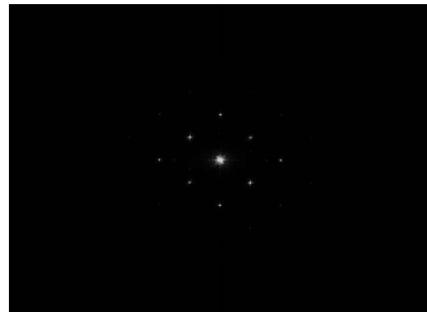
- Periodic noise reduction by selective filters
 - Notch filters can efficiently remove the periodic noise
 - Step 1: analyze the Fourier spectrum F of the image



https://docs.opencv.org/3.4/d2/d0b/tutorial_periodic_noise_removing_filter.html

Periodic Noise Removal: Frequency Domain Filtering

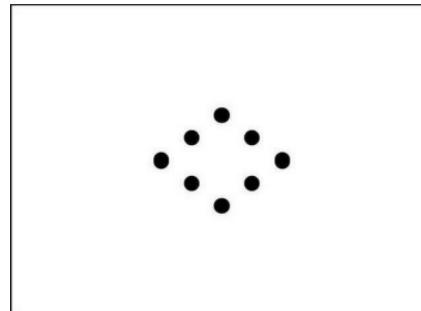
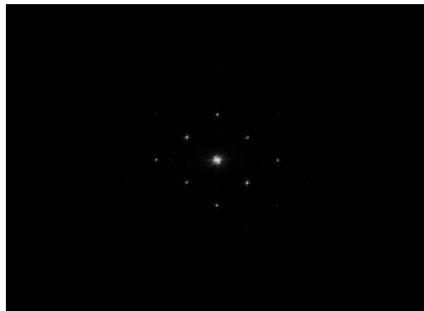
- Periodic noise reduction by selective filters
 - Notch filters can efficiently remove the periodic noise
 - Step 1: analyze the Fourier spectrum F of the image
 - Step 2: identify the locations of the peaks in F



https://docs.opencv.org/3.4/d2/d0b/tutorial_periodic_noise_removing_filter.html

Periodic Noise Removal: Frequency Domain Filtering

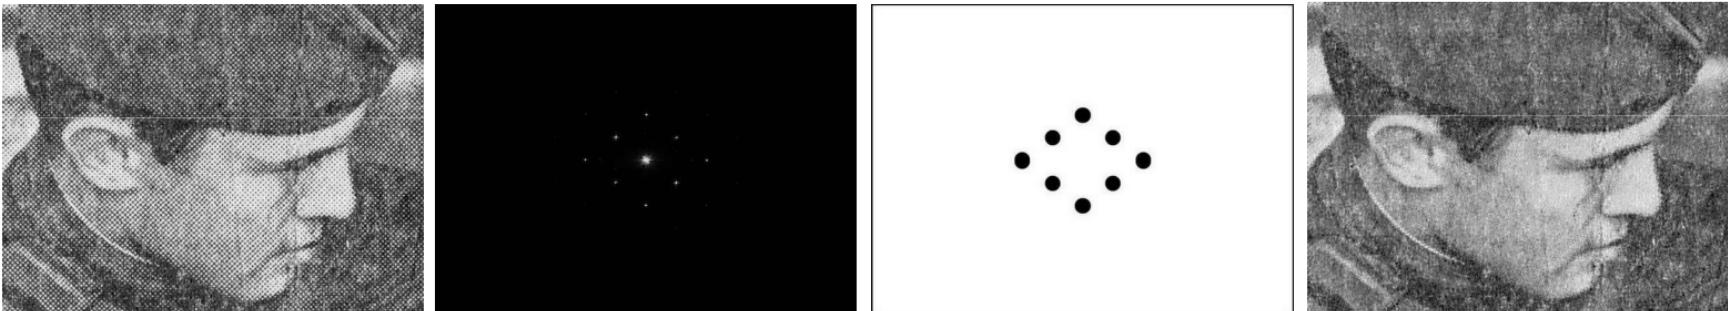
- Periodic noise reduction by selective filters
 - Notch filters can efficiently remove the periodic noise
 - Step 1: analyze the Fourier spectrum F of the image
 - Step 2: identify the locations of the peaks in F
 - Step 3: construct a notch reject filter H in Fourier domain, whose centers are at peaks



https://docs.opencv.org/3.4/d2/d0b/tutorial_periodic_noise_removing_filter.html

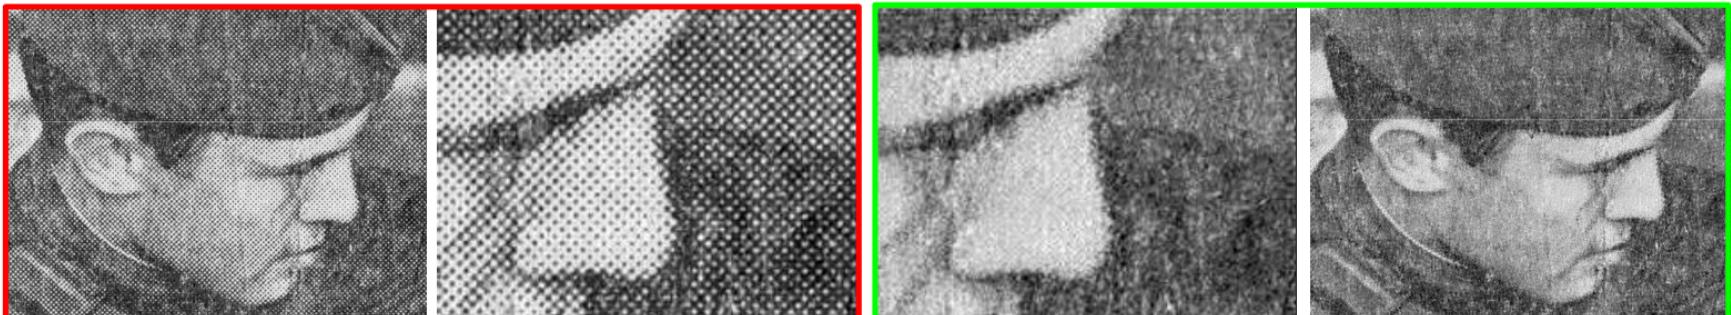
Periodic Noise Removal: Frequency Domain Filtering

- Periodic noise reduction by selective filters
 - Notch filters can efficiently remove the periodic noise
 - Step 1: analyze the Fourier spectrum F of the image
 - Step 2: identify the locations of the peaks in F
 - Step 3: construct a notch reject filter H in Fourier domain, whose centers are at peaks
 - Step 4: use H to filter F to get the result



Periodic Noise Removal: Frequency Domain Filtering

- Periodic noise reduction by selective filters
 - Notch filters can efficiently remove the periodic noise
 - Step 1: analyze the Fourier spectrum F of the image
 - Step 2: identify the locations of the peaks in F
 - Step 3: construct a notch reject filter H in Fourier domain, whose centers are at peaks
 - Step 4: use H to filter F to get the result



https://docs.opencv.org/3.4/d2/d0b/tutorial_periodic_noise_removing_filter.html

Summary

- **4 different type of noises**
 - Salt and Pepper noise, Gaussian noise, Speckle noise, Periodic noise
- **Salt-and-Pepper Noise Removal**
 - Median filtering produces the best results.
- **Gaussian Noise Removal**
 - Additive White Gaussian Noise (AWGN): the most widely used noise model
 - Adaptive filtering produces satisfactory results.
- **Periodic Noise Removal**
 - Filtering in Frequency domain

outlier rejection ↑
(detective whether noise?)

{
Gaussian filter / average ...
— smooth the structure
Bilateral filter (edge)
— artifacts
WNS (non-local means)

Next topic

- **How can we represent images more effectively?**
 - Prerequisite
 - Review EBU6230 Image/Video Processing – Week3: Edges
 - Review EBU6230 Image/Video Processing – Week3: Interest Points

$$S = \{s_1, \dots, s_q\} = \{0.0, 2.2, 4, 6.6, 6, 100\}$$

1) Find x which minimises $y = \sum_{i=1}^q (x - s_i)^2 \Rightarrow L_2$ error
 L₂ minimisation

$$y' = 0 \Rightarrow \sum_{i=1}^q (x - s_i) = 0$$

SSD : sum of squared difference
 (mean value)

$$\sum_{i=1}^q x = \sum_{i=1}^q s_i$$

$$qx = \sum_{i=1}^q s_i \Rightarrow x = \frac{1}{q} \sum_{i=1}^q s_i$$

$x \Rightarrow$ mean value of $s_i \sim s_q$

2) Find x which minimises $y = \sum_{i=1}^q |x - s_i| \Rightarrow L_1$ error

L₁ minimisation

$$y' = \sum_{i=1}^q \text{sgn}(x - s_i)$$

SAD : sum of absolute difference
 (median value)

$$\text{sgn}(t) = \begin{cases} 1 & t > 0 \\ 0 & t = 0 \\ -1 & t < 0 \end{cases}$$

$$-1 -1 -1 0 1 1 1 1$$

$$\uparrow$$

$x = 4$ (median value)

$x \Rightarrow$ median value of $s_i \sim s_q$